



# THE CUDA PLATFORM

Paul Graham, Senior Solutions Architect, NVIDIA

# AGENDA

- Introduction to GPU computing
- CUDA platform overview
- Software architecture trends
- CUDA tools and libraries
- Next steps







# INTRODUCTION TO GPU COMPUTING

# Add GPUs: Accelerate Applications

**CPU**



**GPU**



# TESLA V100 TENSOR CORE GPU

World's Most Powerful  
Data Center GPU

5,120 CUDA cores  
640 NEW Tensor cores  
7.8 FP64 TFLOPS | 15.7 FP32 TFLOPS  
| 125 Tensor TFLOPS  
32 GB HBM2 @ 900GB/s |  
300GB/s NVLink









# NEW TURING TENSOR CORE

MULTI-PRECISION FOR AI TRAINING AND INFERENCE


65 TFLOPS FP16 | 130 TeraOPS INT8 | 260 TeraOPS INT4




# GPU COMPUTING AT THE HEART OF AI

New Advancements Leapfrog Moore's Law

## Performance Beyond Moore's Law




## Big Bang of Modern AI



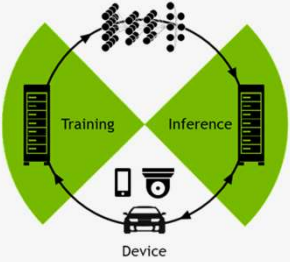
# MOST ADOPTED PLATFORM FOR ACCELERATING AI

8 MLPerf 0.6 Training Records



	Benchmark	Record
At Scale Record	Object Detection (Heavy Weight) Mask R-CNN	18.47 Mins
	Translation (Recurrent) GNMT	1.8 Mins
	Reinforcement Learning (MiniGo)	13.57 Mins
Per Accelerator Record	Object Detection (Heavy Weight) Mask R-CNN	25.39 Hrs
	Object Detection (Light Weight) SSD	3.04 Hrs
	Translation (Recurrent) GNMT	2.63 Hrs
	Translation (Non-recurrent)Transformer	2.61 Hrs
	Reinforcement Learning (MiniGo)	3.65 Hrs

RECORD-SETTING  
PERFORAMNCE



Chainer

mxnet

PYTORCH

TensorFlow

TensorRT

ONNX

python

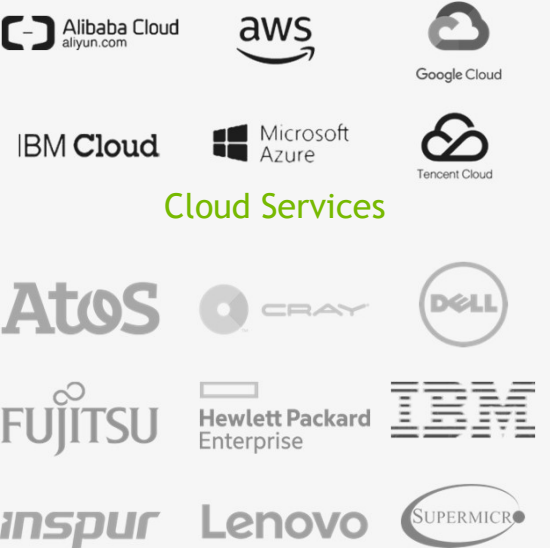
RAPIDS

AWS SageMaker

GCP ML Engine

AzureML

END-TO-END  
SOFTWARE STACK



Cloud Services

Systems

AVAILABLE  
EVERYWHERE



# 3 WAYS TO ACCELERATE APPLICATIONS

Applications

Libraries

“Drop-in”  
Acceleration


OpenACC  
Directives

Easily Accelerate  
Applications

Programming  
Languages

Maximum  
Flexibility

# SMALL CHANGES, BIG SPEED-UP



# 3 WAYS TO ACCELERATE APPLICATIONS

Applications

Libraries

“Drop-in”  
Acceleration

OpenACC  
Directives

Easily Accelerate  
Applications

Programming  
Languages

Maximum  
Flexibility

# LIBRARIES: EASY, HIGH-QUALITY ACCELERATION

**EASE OF USE** Using libraries enables GPU acceleration without in-depth knowledge of GPU programming

**“DROP-IN”** Many GPU-accelerated libraries follow standard APIs, thus enabling acceleration with minimal code changes

**QUALITY** Libraries offer high-quality implementations of functions encountered in a broad range of applications


**PERFORMANCE** NVIDIA libraries are tuned by experts




# GPU ACCELERATED LIBRARIES

“Drop-in” Acceleration for Your Applications


## DEEP LEARNING





## SIGNAL & IMAGE PROCESSING



## LINEAR ALGEBRA



## PARALLEL ALGORITHMS



# 3 STEPS TO CUDA-ACCELERATED APPLICATION

**Step 1:** Substitute library calls with equivalent CUDA library calls

```
saxpy ( ... ) ➤ cublasSaxpy ( ... )
```

**Step 2:** Manage data locality

- with CUDA: `cudaMalloc()`, `cudaMemcpy()`, etc.
- with CUBLAS: `cublasAlloc()`, `cublasSetVector()`, etc.

**Step 3:** Rebuild and link the CUDA-accelerated library

```
gcc myobj.o -l cublas
```

# DROP-IN ACCELERATION (STEP 1)

```
int N = 1 << 20;
```

```
// Perform SAXPY on 1M elements: y[]=a*x[]+y[]  
saxpy(N, 2.0, d_x, 1, d_y, 1);
```

# DROP-IN ACCELERATION (STEP 1)

```
int N = 1 << 20;
```

```
// Perform SAXPY on 1M elements: d_y[]=a*d_x[]+d_y[]  
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);
```

Add “cublas” prefix  
and use device variables



## DROP-IN ACCELERATION (STEP 2)

```
int N = 1 << 20;  
cublasInit();
```

Initialize cuBLAS

```
// Perform SAXPY on 1M elements: d_y[]=a*d_x[]+d_y[]  
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);
```

```
cublasShutdown();
```

Shut down cuBLAS

## DROP-IN ACCELERATION (STEP 2)

```
int N = 1 << 20;
cublasInit();
cublasAlloc(N, sizeof(float), (void**) &d_x);
cublasAlloc(N, sizeof(float), (void**) &d_y);

// Perform SAXPY on 1M elements: d_y[] = a*d_x[] + d_y[]
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);

cublasFree(d_x);
cublasFree(d_y);
cublasShutdown();
```

Allocate device vectors

Deallocate device vectors

## DROP-IN ACCELERATION (STEP 2)

```
int N = 1 << 20;
cublasInit();
cublasAlloc(N, sizeof(float), (void**)&d_x);
cublasAlloc(N, sizeof(float), (void**)&d_y);

cublasSetVector(N, sizeof(x[0]), x, 1, d_x, 1);
cublasSetVector(N, sizeof(y[0]), y, 1, d_y, 1);

// Perform SAXPY on 1M elements: d_y[] = a*d_x[] + d_y[]
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);

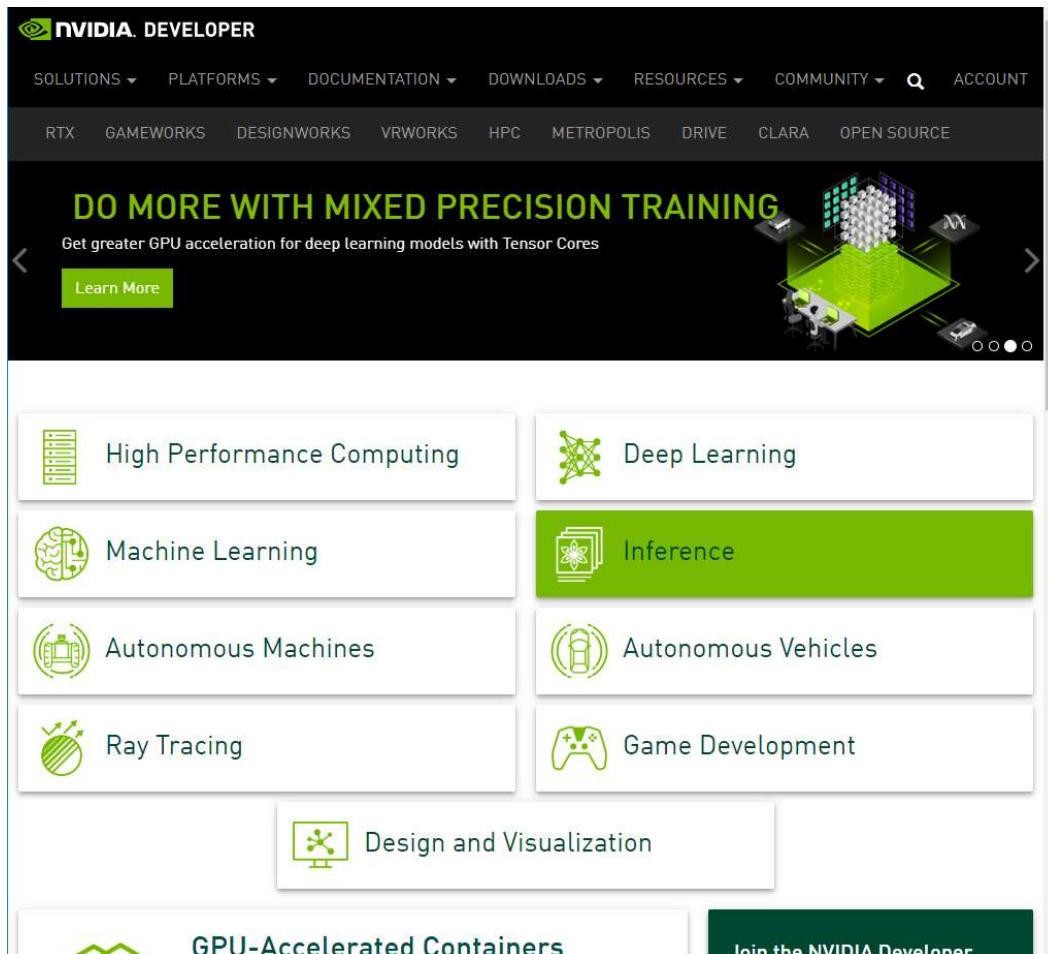
cublasGetVector(N, sizeof(y[0]), d_y, 1, y, 1);

cublasFree(d_x);
cublasFree(d_y);
cublasShutdown();
```

Transfer data to GPU

Read data back GPU

# EXPLORE CUDA LIBRARIES



[developer.nvidia.com](https://developer.nvidia.com)



# 3 WAYS TO ACCELERATE APPLICATIONS

Applications

Libraries

“Drop-in”  
Acceleration

OpenACC  
Directives

Easily Accelerate  
Applications

Programming  
Languages

Maximum  
Flexibility

**OpenACC** is a directives-based programming approach to **parallel computing** designed for **performance** and **portability** on CPUs and GPUs for HPC.


Add Simple Compiler Directive

```
main()
{
    <serial code>
    #pragma acc kernels
    {
        <parallel code>
    }
}
```



# LSDALTON

Large-scale application for calculating high-accuracy molecular energies



“OpenACC makes GPU computing approachable for domain scientists. Initial OpenACC implementation required only *minor effort, and more importantly, no modifications* of our existing CPU implementation.”

Janus Juul Eriksen, PhD Fellow  
qLEAP Center for Theoretical Chemistry, Aarhus University



”


## Minimal Effort

Lines of Code Modified	# of Weeks Required	# of Codes to Maintain
<100 Lines	1 Week	1 Source

## Big Performance

LS-DALTON CCSD(T) Module

Benchmarked on Titan Supercomputer (AMD CPU vs Tesla K20X)




<https://developer.nvidia.com/openacc/success-stories>

# SINGLE CODE FOR MULTIPLE PLATFORMS

## OpenACC - Performance Portable Programming Model for HPC

POWER  
Sunway  
x86 CPU  
x86 Xeon Phi  
NVIDIA GPU  
PEZY-SC



Systems: Haswell: 2x16 core Haswell server, four K80s, CentOS 7.2 (perf-hsw10), Broadwell: 2x20 core Broadwell server, eight P100s (dgx1-prd-01), Minsky: POWER8+NVLINK, four P100s, RHEL 7.3 (gsn1).  
Compilers: Intel 17.0, IBM XL 13.1.3, PGI 16.10.  
Benchmark: CloverLeaf v1.3 downloaded from <http://uk-mac.github.io/CloverLeaf> the week of November 7 2016; CloverLeaf\_Serial; CloverLeaf\_ref (MPI+OpenMP); CloverLeaf\_OpenACC (MPI+OpenACC)  
Data compiled by PGI November 2016, Volta data collected June 2017



## 2 BASIC STEPS TO GET STARTED

### Step 1:

```
!$acc data copy(util1,util2,util3) copyin(ip,scp2,scp2i)
!$acc parallel loop
...
!$acc end parallel
!$acc end data
```

### Step 2:

```
pgf90 -ta=nvidia -Minfo=accel file.f
```

# OpenACC DIRECTIVES EXAMPLE

```
!$acc data copy(A,Anew)
```

```
iter=0  
do while ( err > tol .and. iter < iter_max )  
  
    iter = iter +1  
    err=0._fp_kind
```

```
!$acc kernels
```

```
    do j=1,m  
        do i=1,n  
            Anew(i,j) = .25_fp_kind *( A(i+1,j ) + A(i-1,j ) &  
                                     +A(i ,j-1) + A(i ,j+1))  
  
            err = max( err, Anew(i,j)-A(i,j))  
        end do  
    end do
```

```
!$acc end kernels
```

```
    IF(mod(iter,100)==0 .or. iter == 1)    print *, iter, err  
    A= Anew
```

```
end do
```

```
!$acc end data
```



Copy arrays into GPU memory  
within data region



Parallelize code inside region



Close off parallel region



Close off data region,  
copy data back

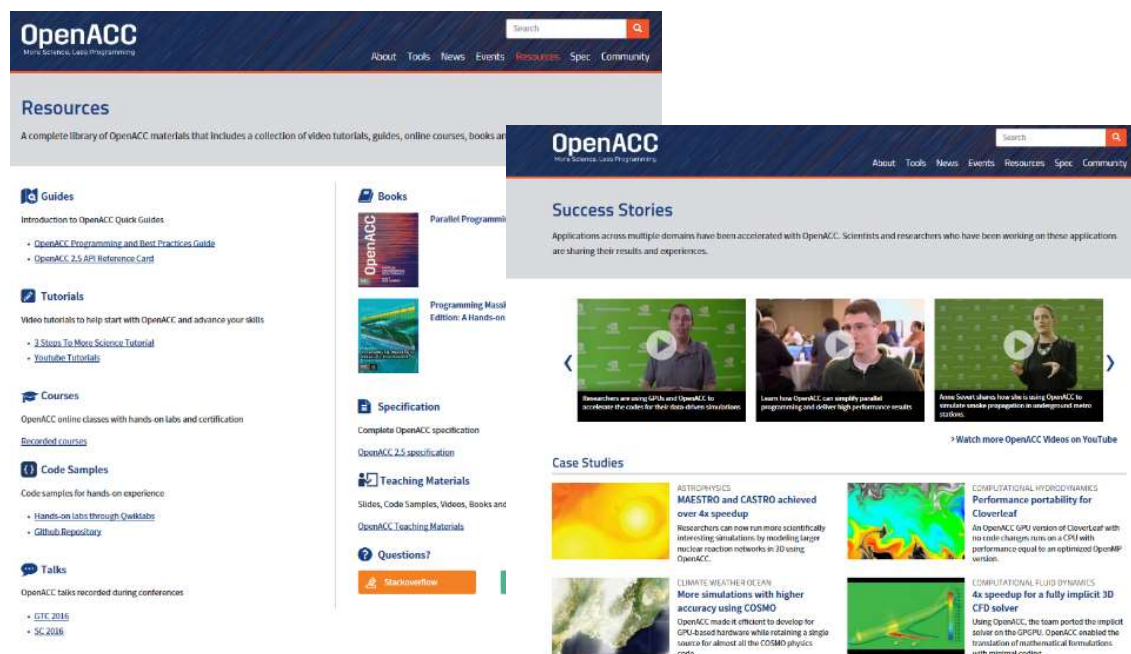
# OPENACC FOR EVERYONE

PGI Community Edition Available

	FREE		
	<b>PGI</b> Community EDITION	<b>PGI</b> Professional EDITION	<b>PGI</b> Enterprise EDITION
<b>PROGRAMMING MODELS</b> OpenACC, CUDA Fortran, OpenMP, C/C++/Fortran Compilers and Tools	✓	✓	✓
<b>PLATFORMS</b> X86, OpenPOWER, NVIDIA GPU	✓	✓	✓
<b>UPDATES</b>	1-2 times a year	6-9 times a year	6-9 times a year
<b>SUPPORT</b>	User Forums	PGI Support	PGI Professional Services
<b>LICENSE</b>	Annual	Perpetual	Volume/Site

# RESOURCES

FREE Compiler  
Success stories  
Guides  
Tutorials  
Videos  
Courses  
Code Samples  
Talks  
Books Specification  
Teaching Materials  
Slack&StackOverflow



Success stories: <https://www.openacc.org/success-stories>

Resources: <https://www.openacc.org/resources>

Free Compiler: <https://www.pgroup.com/products/community.htm>

# 3 WAYS TO ACCELERATE APPLICATIONS

Applications

Libraries

“Drop-in”  
Acceleration

OpenACC  
Directives

Easily Accelerate  
Applications

Programming  
Languages

Maximum  
Flexibility

# GPU PROGRAMMING LANGUAGES

Numerical analytics ►

MATLAB, Mathematica, LabVIEW

Fortran ►

CUDA Fortran, OpenACC

C, C++ ►

CUDA C++, OpenACC

Python ►

CUDA Python, PyCUDA

C# ►

Altimesh Hybridizer, Alea GPU

# CUDA C

```
void saxpy_serial(int n,
                  float a,
                  float *x,
                  float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

// Perform SAXPY on 1M elements
saxpy_serial(4096*256, 2.0, x, y);
```

```
__global__
void saxpy_parallel(int n,
                    float a,
                    float *x,
                    float *y)
{
    int i = blockIdx.x*blockDim.x +
            threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

// Perform SAXPY on 1M elements
saxpy_parallel<<<4096,256>>>(n,2.0,x,y);
```

<http://developer.nvidia.com/cuda-toolkit>



# CUDA C++: DEVELOP GENERIC PARALLEL CODE

CUDA C++ features enable sophisticated and flexible applications and middleware

Class hierarchies

\_\_device\_\_ methods

Templates

Operator overloading

Functors (function objects)

Device-side new/delete

More...

<http://developer.nvidia.com/cuda-toolkit>

```
template <typename T>
struct Functor {
    __device__ Functor(_a) : a(_a) {}
    __device__ T operator(T x) { return a*x; }
    T a;
}

template <typename T, typename Oper>
__global__ void kernel(T *output, int n) {
    Oper op(3.7);
    output = new T[n]; // dynamic allocation
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n)
        output[i] = op(i); // apply functor
}
```

# RAPID PARALLEL C++ DEVELOPMENT

- Resembles C++ STL
- High-level interface
  - Enhances developer productivity
  - Enables performance portability between GPUs and multicore CPUs
- Flexible
  - CUDA, OpenMP, and TBB backends
  - Extensible and customizable
  - Integrates with existing software
- Open source



```
// generate 32M random numbers on host
thrust::host_vector<int> h_vec(32 << 20);
thrust::generate(h_vec.begin(),
                 h_vec.end(),
                 rand);

// transfer data to device (GPU)
thrust::device_vector<int> d_vec = h_vec;
// sort data on device
thrust::sort(d_vec.begin(), d_vec.end());
// transfer data back to host
thrust::copy(d_vec.begin(),
             d_vec.end(),
             h_vec.begin());
```

<http://developer.nvidia.com/thrust> or <http://thrust.googlecode.com>

# CUDA FORTRAN

- Program GPU using Fortran
  - Key language for HPC
- Simple language extensions
  - Kernel functions
  - Thread / block IDs
  - Device & data management
  - Parallel loop directives
- Familiar syntax
  - Use allocate, deallocate
  - Copy CPU-to-GPU with assignment (=)

<http://developer.nvidia.com/cuda-fortran>

```
module mymodule contains
  attributes(global) subroutine saxpy(n,a,x,y)
    real :: x(:), y(:), a,
    integer n, i
    attributes(value) :: a, n
    i = threadIdx%x+(blockIdx%x-1)*blockDim%x
    if (i<=n) y(i) = a*x(i) + y(i);
  end subroutine saxpy
end module mymodule

program main
  use cudafor; use mymodule
  real, device :: x_d(2**20), y_d(2**20)
  x_d = 1.0; y_d = 2.0
  call saxpy<<<4096,256>>>(2**20,3.0,x_d,y_d,)
  y = y_d
  write(*,*) 'max error=', maxval(abs(y-5.0))
end program main
```

# GET STARTED TODAY

These languages are supported on all CUDA-capable GPUs.

You might already have a CUDA-capable GPU in your laptop or desktop PC!

CUDA C/C++

<http://developer.nvidia.com/cuda-toolkit>

CUDA Python

<http://developer.nvidia.com/how-to-cuda-python>

Thrust C++ Template Library

<http://developer.nvidia.com/thrust>

CUDA Fortran

<http://developer.nvidia.com/cuda-toolkit>

MATLAB

<http://www.mathworks.com/discovery/matlab-gpu.html>

Mathematica

<http://www.wolfram.com/mathematica/new-in-8/cuda-and-opencl-support/>



# **SIX WAYS TO SAXPY**

Programming Languages for GPU Computing

# SINGLE PRECISION ALPHA X PLUS Y (SAXPY)

Part of Basic Linear Algebra Subroutines (BLAS) Library

$$z = \alpha x + y$$

$x, y, z$  : vector

$\alpha$  : scalar

GPU SAXPY in multiple languages and libraries

A menagerie\* of possibilities, not a tutorial

\*technically, a *program chrestomathy*: <http://en.wikipedia.org/wiki/Chrestomathy>

1

# OpenACC COMPILER DIRECTIVES

## *Parallel C Code*

```
void saxpy(int n,
           float a,
           float *x,
           float *y)
{
    #pragma acc kernels
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

...
// Perform SAXPY on 1M elements
saxpy(1<<20, 2.0, x, y);
...
```

## *Parallel Fortran Code*

```
subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    !$acc kernels
    do i=1,n
        y(i) = a*x(i)+y(i)
    enddo
    !$acc end kernels
end subroutine saxpy

...
! Perform SAXPY on 1M elements
call saxpy(2**20, 2.0, x_d, y_d)
...
```

<http://developer.nvidia.com/openacc> or <http://openacc.org>

## 2

# cuBLAS LIBRARY

## *Serial BLAS Code*

```
int N = 1<<20;

...

// Use your choice of blas library

// Perform SAXPY on 1M elements
blas_saxpy(N, 2.0, x, 1, y, 1);
```

## *Parallel cuBLAS Code*

```
int N = 1<<20;
cublasInit();
cublasAlloc(N, sizeof(float), (void**)&d_x);
cublasAlloc(N, sizeof(float), (void**)&d_y);
cublasSetVector(N, sizeof(x[0]), x, 1, d_x, 1);
cublasSetVector(N, sizeof(y[0]), y, 1, d_y, 1);

// Perform SAXPY on 1M elements
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);

cublasGetVector(N, sizeof(y[0]), d_y, 1, y, 1);

cublasShutdown();
```

You can also call cuBLAS from Fortran,  
C++, Python, and other languages

<http://developer.nvidia.com/cublas>



3

## CUDA C

```
void saxpy(int n, float a,
          float *x, float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}
```

```
int N = 1<<20;
```

```
// Perform SAXPY on 1M elements
saxpy(N, 2.0, x, y);
```

```
__global__
void saxpy(int n, float a,
          float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}
```

```
int N = 1<<20;
cudaMemcpy(d_x, x, N, cudaMemcpyHostToDevice);
cudaMemcpy(d_y, y, N, cudaMemcpyHostToDevice);
```

```
// Perform SAXPY on 1M elements
saxpy<<<4096,256>>>(N, 2.0, d_x, d_y);

cudaMemcpy(y, d_y, N, cudaMemcpyDeviceToHost);
```

<http://developer.nvidia.com/cuda-toolkit>

## 4

# THRUST C++ TEMPLATE LIBRARY

## *Serial C++ Code with STL and Boost*

```
int N = 1<<20;
std::vector<float> x(N), y(N);

...

// Perform SAXPY on 1M elements
std::transform(x.begin(), x.end(),
               y.begin(), y.end(),
               2.0f * _1 + _2);
```

[www.boost.org/libs/lambda](http://www.boost.org/libs/lambda)

## *Parallel C++ Code*

```
int N = 1<<20;
thrust::host_vector<float> x(N), y(N);

...

thrust::device_vector<float> d_x = x;
thrust::device_vector<float> d_y = y;

// Perform SAXPY on 1M elements
thrust::transform(d_x.begin(), d_x.end(),
                  d_y.begin(), d_y.begin(),
                  2.0f * _1 + _2);
```

<http://thrust.github.com>

# CUDA FORTRAN

## Standard Fortran

```

module mymodule contains
  subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    do i=1,n
      y(i) = a*x(i)+y(i)
    enddo
  end subroutine saxpy
end module mymodule

program main
  use mymodule
  real :: x(2**20), y(2**20)
  x = 1.0, y = 2.0
  ! Perform SAXPY on 1M elements
  call saxpy(2**20, 2.0, x, y)
end program main

```

## Parallel Fortran

```

module mymodule contains
  attributes(global) subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    attributes(value) :: a, n
    i = threadIdx%x+(blockIdx%x-1)*blockDim%x
    if (i<=n) y(i) = a*x(i)+y(i)
  end subroutine saxpy
end module mymodule

program main
  use cudafor; use mymodule
  real, device :: x_d(2**20), y_d(2**20)
  x_d = 1.0, y_d = 2.0
  ! Perform SAXPY on 1M elements
  call saxpy<<<4096,256>>>(2**20, 2.0, x_d, y_d)
end program main

```

<http://developer.nvidia.com/cuda-fortran>

# PYTHON

## *Standard Python*

```
import numpy as np

def saxpy(a, x, y):
    return [a * xi + yi
            for xi, yi in zip(x, y)]

x = np.arange(2**20, dtype=np.float32)
y = np.arange(2**20, dtype=np.float32)

cpu_result = saxpy(2.0, x, y)
```

<http://numpy.scipy.org>

## *Numba Parallel Python*

```
import numpy as np
from numba import vectorize

@vectorize(['float32(float32, float32, float32)',
           'float32(float32, float32, float32)'], target='cuda')
def saxpy(a, x, y):
    return a * x + y

N = 1048576

# Initialize arrays
A = np.ones(N, dtype=np.float32)
B = np.ones(A.shape, dtype=A.dtype)
C = np.empty_like(A, dtype=A.dtype)

# Add arrays onGPU
C = saxpy(2.0, X, Y)
```

<https://numba.pydata.org>

# ENABLING ENDLESS WAYS TO SAXPY

- Build front-ends for Java, Python, R, DSLs
- Target other processors like ARM, FPGA, GPUs, x86

**CUDA Compiler Contributed to  
Open Source LLVM**

